

Naijith Thandilakath

Software Engineer | AI/ML Engineer

gnaijith78@gmail.com ♦ [+91 9497621395](tel:+919497621395) ♦ linkedin.com/in/naijith ♦ naijith.dev ♦ github.com/naijith

Bangalore, India

SUMMARY

Software and AI/ML engineer with 6+ years of experience building systems that go to production and stay there: microservices, agentic AI applications, and cloud-native platforms. Most comfortable working across the full lifecycle from architecture to deployment.

SKILLS

Languages	Java, Python, SQL, Golang
Backend	Spring Boot, FastAPI, Microservices, REST, gRPC, Protobuf, Temporal
AI/ML & Data	LangGraph, Azure AI Foundry, Agentic AI, RAG, PySpark, Delta Lake, Kafka, CDC
Databases	PostgreSQL, MongoDB, Redis, Azure Cosmos DB
Cloud & DevOps	Azure, AWS, GCP, Docker, Kubernetes, Helm, Terraform, Git, CI/CD, GitHub Actions

EXPERIENCE

Software and AI/ML Engineer

Oct 2024 – Present

A.P. Moller - Maersk

Bangalore, India

- Architected Rate Management System and RateFlow (Java, Spring Boot, Temporal, PostgreSQL) using DDD and hexagonal architecture, processing 500K+ transactions/month across 7 microservices for ocean rate audit and order creation.
- Engineered Rosetta, agentic document parser (Python, FastAPI, LangGraph) that auto-generates vendor-specific parsers via LLMs and self-corrects on low-confidence outputs, achieving 90% accuracy across PDFs, Excel, and images.
- Designed custom query engine and domain-specific query language in Python, backed by PySpark and Delta Lake, replacing manual SQL workflows by data engineers with a self-serve UI that enables business users to run ad-hoc analytics over 2 billion+ IoT records.
- Drove Kafka event streaming with CDC connectors, replacing nightly batch syncs with real-time propagation of rate and order changes, eliminating DB throttling.
- Optimized PostgreSQL and Cosmos DB schemas for high-throughput workloads, maintaining p95 API latency under 500ms.
- Delivered cross-regional deployment APIs for MDP (Maersk Developer Platform), enabling multi-region Kubernetes rollouts with HA/DR and zero-downtime releases.

Software Engineer

Oct 2022 – Oct 2024

IBM Software Labs

Kochi, India

- Reduced monthly processing cost by 80% by migrating workers from paid SPS pipelines to Kubernetes-orchestrated containers.
- Accelerated security scan and network-validation services for Db2 and Watsonx.data, cutting scanning time by 70% through Kafka partitioning and horizontally scaled workers.
- Shipped cloud provisioning APIs in Golang using gRPC/Protobuf for Watsonx.data, maintained Helm charts for highly available multi-cloud deployments.
- Built analytics APIs (FastAPI, PostgreSQL) with 500ms latency, powering time-series dashboards that tracked vulnerability trends across releases, enabling teams to measure remediation progress and spot regressions early.
- Instrumented centralized observability using ELK and Prometheus for microservices logging and monitoring.

Software Engineer

TCS (Marks & Spencer)

Sep 2020 – Oct 2022

Kochi, India

- Migrated APIs from on-prem Kubernetes to Azure AKS, cutting infrastructure costs by 30% and operational overhead by 25%.
- Owned Java Spring Boot microservices for ETL processing of IoT-driven product and inventory data, storing enriched datasets in MongoDB for downstream retail systems.
- Launched inventory reconciliation service validating ETL accuracy, enabling early detection of data drift and load failures.
- Standardized API management using Apigee with centralized authentication, policy enforcement, and traffic controls.
- Automated New Relic alerting with Terraform and PagerDuty integration, reducing MTTR by 40%.

PROJECTS

Breaking the Lines – *Football fan community platform*

breakingthelines.com

- Co-built entire platform from scratch in a 2-person team (Golang, PostgreSQL, AWS S3, Redis), shipping 3 core services over a Consul/Envoy service mesh with Centrifugo for real-time match updates and notifications, now serving 1M+ hits/month.

AI Portfolio Chatbot

naijith.dev

- Built interactive portfolio with client-side RAG chat – loads a quantized MiniLM-L6 model (q8) via Hugging Face Transformers.js, embeds queries in-browser, and retrieves context from pre-computed embeddings using cosine similarity, with zero backend or API costs.

EDUCATION

Bachelor of Computer Science

SCMS School of Engineering and Technology

2016 – 2020

Kochi, India